

ALTeGraD 2023 Data Challenge

Molecule Retrieval with Natural Language Queries

L. Hunout S. Hocine L. Haubert

Master MVA
École Normale Supérieure Paris-Saclay

February 14, 2024

Table of Contents

- 1 Task Definition
- 2 Methodology
- 3 Experiments
- 4 Conclusion and Future Work
- 5 Appendices

An example of molecule retrieval

Natural language query → *Structure of the related molecule*

Cholesteryl linoleate is the (9Z,12Z)-stereoisomer of cholesteryl octadeca-9,12-dienoate. It has a role as a human metabolite and a mouse metabolite. It derives from a linoleic acid.




Figure: The idea of molecule retrieval with natural language queries

About the mission

■ Context:

- In sciences, medicine, biology, it may be needed to exploit text-molecule relations, e.g for retrieval tasks
- Retrieval involves obtaining relevant data from a user query

■ Definition:

- Extract molecules based on natural language descriptions
- Queries are textual data and molecules are graph data

■ Key considerations:

- Integrating natural language and molecules is complex, due to their different information encoding
- The challenge consists in building a relevant comparison tool to settle relevant matching scores

Solving and evaluation paradigm

■ Contrastive Learning (CL):

- Text and graph encoders are jointly trained with CL
- CL maps the similar text-graph pairs close together in the representation space and push other ones apart
- Contrastive loss is defined with the cross-entropy

■ Label Ranking Average Precision score (LRAP):

- Evaluation metric to quantify the relevance of the retrieval
- Only one molecule to retrieve \implies LRAP \sim MRR
- MRR: Average, over the queries, of the inverted retrieval ranks

■ About what follows:

- Same paradigm for the baseline and the contributions
- What will change: models, trainings, adaptations...

Baseline model

- **What was given as a source code:**
 - A pipeline was provided to give a simple baseline
 - LRAP score on half of the test data: about 30%
- **Details about the baseline:**
 - Text encoder: DistilBERT
 - Graph encoder: GCN (Graph Convolutional Network)
 - Training: 5 epochs, size of batch 32, learning rate 2e-5
- **Our mission:**
 - Develop a new retrieval approach to obtain better LRAP scores
 - Analyze and use, if needed, other text and graph encoders
 - Leverage optimization techniques to improve the pipeline

Contrastive Learning Loss

- Objective: Minimize distance between similar embeddings and maximize for dissimilar pairs
- Cross Entropy (CE) loss function is used for this purpose
- For each graph-text pair, the loss aims to:
 - Maximize probability of correct matching
 - Minimize probability of incorrect matching
- The formula for the contrastive loss using CE is:

$$\mathcal{L}_{CL} = -\frac{1}{N} \sum_{i=1}^N \left[\log \left(\frac{\exp(s(z_{\text{graph}_i}, z_{\text{text}_i}))}{\sum_{j=1}^N \exp(s(z_{\text{graph}_i}, z_{\text{text}_j}))} \right) + \log \left(\frac{\exp(s(z_{\text{text}_i}, z_{\text{graph}_i}))}{\sum_{j=1}^N \exp(s(z_{\text{text}_i}, z_{\text{graph}_j}))} \right) \right]$$

InfoNCE Loss Exploration

- Initially explored InfoNCE loss due to its prevalence in contrastive learning literature
- InfoNCE loss formula:

$$\mathcal{L}_{\text{InfoNCE}} = -\log \frac{\exp(s(z_i, z_j)/\tau)}{\sum_{k=1}^N \exp(s(z_i, z_k)/\tau)}$$

- Despite it's widely used in contrastive learning, it led to poor convergence in our experiments

Text Encoder

■ DistilBERT **model**:

- Model provided in the baseline, initially adopted
- Efficient and light due to the distillation process from BERT
- Train a smaller model to replicate a larger one (here BERT)

■ SciBERT **model**:

- Other tested model, to refine our approach towards efficiency
- BERT-based architecture, pre-trained on scientific texts

■ **Comparison of the two models**:

- SciBERT: Better convergence rate (which was expected)
- Both models achieved similar scores/losses upon convergence
- DistilBERT: Light architecture \implies larger batch size and lower complexity (training) \implies practical advantage
- We keep DistilBERT as the text encoder

Graph Encoders

- **Common goal:** Encode graph info into numerical representations
- **Shared characteristics:**
 - Designed for graph data
 - Input: Node features, connections. Output: Dense vector
 - Convolutional layers consider graph structure
 - Regularization: Batch Normalization, Dropout
- **Differences:**
 - Types of convolutional layers used
 - Operation on graph data

Graph Encoders: Differences

1 GCNEncoder:

- Uses GCNConv layers to convolve node features, capturing local relationships
- Aggregates neighbor features without differentiation

2 GATEncoder:

- Utilizes GATv2Conv layers with attention mechanisms, focusing on specific parts of the graph
- More effective in contexts where relationships between nodes are heterogeneous and the relative importance of neighboring nodes varies

3 GINEncoder:

- Distinguish non isomorphic graphs, relying on the Weisfeiler-Lehman isomorphism test
- Flexible in representation, doesn't presume structure
- Enable to get a fine understanding of graph topology

Ensemble Approach

- Correct molecule frequently found among top-ranked molecules
- Matches found jointly by all base models, while some found by certain models only due to differing feature extraction by graph encoders
- Objective: Utilize capabilities of different models via ensemble approach
- Approach: Average results obtained for each model (n-to-n cosine similarity matrix)
- Noticed: Normalizing matrices on rows had little impact on operation result

Optimizations for Training Process (1/2)

- Technical limitations encountered:
 - Training on a single Tesla P100 PCIe 16 GB
 - Time restriction of maximum 30 hours per week
- Realization of the value of time for training models efficiently
- Implemented optimizations to overcome challenges:
 - Optimized DataLoader using Torch library for efficient data loading and reduced wait times
 - Incorporated autocast mixed-precision for exploiting 16-bit floating point calculations, providing significant speedup without compromising results quality
 - Adopted Hugging Face accelerator with Hugging Face library to fully leverage modern hardware capabilities for accelerated training

Optimizations for Training Process (2/2)

- Significance of optimizations:
 - Crucial in addressing CUDA memory errors when processing large volumes of data
 - Enabled increased batch sizes, optimized code, and achieved faster training times
- Contribution to better overall model efficiency

Data and Evaluation

- Data Challenge on Kaggle provides:
 - Host environment and text-graph data
 - Files: train.tsv, val.tsv, token_embedding_dict.npy, cid.graph
 - Test files: test_text.txt, test_cids.txt
- Datasets split into train(80%)/val(10%)/test(10%)
- Shuffled data during training
- Models trained on train data, evaluated by searching all molecules

Results: Baseline Models

- Training details:
 - Adam optimizer with two learning rates:
 - DistilBERT and SciBERT: $3e-5$
 - Others: $1e-4$
 - StepLR scheduler reduces LR by 50% every 5 epochs
 - Trained for 40 epochs with batch size 64
 - First 256 text tokens used for text encoder
 - Batch size upgraded to 128 with DistilBERT

Results: Baseline Models

■ Baseline models:

- GCN, GAT, Stacked-GAT, and GIN show similar performance
- Slight edge for Stacked-GAT and GIN with more weight
- Similar performance attributed to description bottleneck

Model	Training LRAP	Validation LRAP
DistilBERT + GCN	54.26%	75.10%
DistilBERT + GAT	54.53%	74.67%
DistilBERT + Stacked-GAT	57.68%	77.67%
DistilBERT + GIN	57.49%	77.02%
All-Ensemble	NaN	85.68%

Table 1: Results

Results: Ensemble

- Observations:
 - Better LRAPs in validation than training due to dataset size difference
 - LRAP score highly dependent on dataset size
 - Freezing first attention modules in DistilBERT did not improve results
- Ensemble method significantly improves performance
- Increases validation LRAP by 9.5% compared to baseline
- Models with same hyperparameters learn complementary classification methods
- Ensemble incorporates diverse encoder architectures, enhancing efficiency

Conclusion and Future Work

- Study concluded with adoption of DistilBERT as text encoder for its practical advantages
- Ensemble method combining GCN, GAT, and GIN models resulted in significant LRAP score improvements, up to 85.68%
- Future work:
 - Explore data augmentation techniques for more efficient training
 - Consider larger batch sizes for enhanced learning
 - Extend training duration to uncover nuanced patterns
 - Adopt cosine scheduler for dynamic learning rate adjustment
 - Reshuffle and re-split combined datasets to minimize bias
- Intersection of textual and molecular data via machine learning offers new avenues for exploration, highlighting continuous learning in AI

Graph Attention Network v2 (GATv2)

- GATv2 enhances GAT by allowing direct interaction of the attention mechanism with transformed features
- Attention formula:
$$\alpha_{ij} = \text{softmax}_j \left(\text{LeakyReLU} \left(\mathbf{a}^T [W\mathbf{h}_i^{(l)} \parallel W\mathbf{h}_j^{(l)}] \right) \right)$$
- Update rule:
$$\mathbf{h}_i^{(l+1)} = \sigma \left(\sum_{j \in \mathcal{N}(i)} \alpha_{ij} \cdot W\mathbf{h}_j^{(l)} \right)$$
 - $\mathcal{N}(i)$ - set of neighbors of node i
 - α_{ij} - attention coefficients
 - W - weight matrix
 - σ - nonlinear activation function
- Increases the expressiveness of the attention mechanism

Graph Isomorphism Networks (GIN)

- GINs aim to match the power of the Weisfeiler-Lehman graph isomorphism test
- Update rule:
$$h_i^{(l+1)} = \text{MLP}^{(l)} \left((1 + \epsilon^{(l)}) \cdot h_i^{(l)} + \sum_{j \in \mathcal{N}(i)} h_j^{(l)} \right)$$
 - $\mathcal{N}(i)$ - set of neighbors
 - $\epsilon^{(l)}$ - learnable parameter or fixed scalar
 - $\text{MLP}^{(l)}$ - multi-layer perceptron at layer l
- Effectively captures the graph structure in node representations

Data Augmentation

- Contrastive learning benefits from large datasets
- Graph-text contrastive learning has less data compared to image-text
- Data augmentation can significantly improve model performance such as in DeCLIP

Data Augmentation

- **Text augmentation using EDA:**
 - Synonym substitution, random word shuffling, and word elimination
 - Generates noisy versions of text descriptions to enhance training data
- **Graph augmentation using GraphCL methods:**
 - Effective methods for biological molecule graphs: node dropping and subgraph generation
 - These techniques create variations of the original graph, enriching the dataset

Augmentation Process and Challenges

- For each pair, generate 2 augmented texts and 2 augmented graphs
- Compute embeddings and calculate loss for all augmented pairs and average them
- **Implementation issue:** Increased GPU memory cost hindered successful implementation